

Creating a Simple Virtual Machine in C

(With paging of course)

Goal:

Today's society is complex and fast and dare I say overwhelming. In times like these it is normal to yearn simpler times. Maybe for this reason our 4th homework is to implement a virtual machine with its own "operating system". The basic functionalities are similar to a 16 bit VM made by [Andrei Ciobanu](#).

The homework is to add paging for memory virtualization. This is needed to have more than one program in our virtual machine, not that far simpler times. The functions to fill were provided, the job of this vm is to virtualize the memory for processes. That is almost all of it.

Helper Functions And Definitions

```
139
140 // YOUR CODE STARTS HERE
141 size_t getFileSize(FILE *file) {
142     fseek(file, 0, SEEK_END); // Move to the end of the file
143     size_t size = ftell(file); // Get the position (file size)
144     rewind(file); // Reset the file pointer to the beginning
145     return size / sizeof(uint16_t); // Return size in terms of 16-bit words
146 }
147 #define VALID_MASK (0x0001)
148 #define WRITE_MASK (0x0004)
149 #define PTE_SHIFT (11)
150 #define PCB_START (12)
151 uint16_t translateAddress(uint16_t va);
152
153
154 uint16_t getPCB(uint16_t pid)
155 {
156     uint16_t pcb = PCB_START + pid * PCB_SIZE;
157     return pcb;
158 }
159 uint16_t getPageTableStart(uint16_t pid)
160 {
161     //they start from 3rd table
162     uint16_t _pageTableStart = PAGE_SIZE;
163     return (_pageTableStart + pid * 32);
164 }
```

getFileSize is a simple function that returns the size of a file by going to the end of the file, getting that position returning the total number of byte in function, resetting the file pointer and then returning how many 16bits the file is by dividing the size by sizeof(uint16_t). This is needed for the load image provided. By knowing the file size it will write to the position other allocMem function chooses.

getPCB function returns the start of Process Control Block that belongs to the specified pid.

getPageTableStart will find the Page Table belonging to pid. The page tables start from the 3rd page and are reserved 32 per process.

```
284 uint16_t getPTE(uint16_t va)
285 {
286     uint16_t pcb_address = getPCB(mem[Cur_Proc_ID]);
287     uint16_t pte_start = mem[pcb_address + PTBR_PCB];
288     uint16_t vpn = va >> PTE_SHIFT; // Extract VPN
289
290     return mem[(pte_start + vpn)];
291 }
292 uint16_t translateAddress(uint16_t va) {
293     uint16_t pcb_address = getPCB(mem[Cur_Proc_ID]);
294     uint16_t pte_start = mem[pcb_address + PTBR_PCB];
295
296     uint16_t vpn = va >> PTE_SHIFT; // Extract VPN
297     uint16_t offset = va & 0x07FF; // Offset within the page
298
299     //printf("ptbr is %d vpn is ")
300
301     //printf("va is %d and pte_start is %d and vpn is %d",va, pte_start, vpn );
302     uint16_t pte = pte_start + vpn;
303
304     if (!(mem[pte] & VALID_MASK)) { // Check if the page is valid
305         //printf("vpn is %d\n", vpn);
306         fprintf(stderr, "Invalid page table entry for pte: 0x%04X\n", pte); //vpn 0dan başlıyo
307         exit(1); // Terminate to prevent further invalid accesses
308     }
309     pte = mem[pte];
310     uint16_t phys_frame = pte >> PTE_SHIFT;
311     //printf("translated number is %d \n", ((phys_frame << PTE_SHIFT) + offset));
312     return (phys_frame << PTE_SHIFT) + offset;
313 }
```

getPTE returns the page table entry. First finds the process control block, then the page table starting point, and returns the page table entry of the virtual address.

translateAddress get a virtual address. Finds the virtual page number and offset inside the page. We find the pte and check if it is valid. If it is we get the physical page number from it. And return the correct physical memory location.

Definitions

In page table entries the least significant bit is the valid bit. Since I have seen many times in slides the use of macro MASKS in these type of low levelish stuff I preferred a definition for it.

Again for writing permission bit is the 3rd from right. Since 0x0004 is 0000 0000 0000 0100 it will be more readable for these checks.

Page table entries keep the page num in first 5 bits. To access them we have to shift to right 11 times. Since it is cooler to have a macro, I opted for it.

Process Control Block starts at 12th word, however it isn't provided in macros. This one is for ease of remembering.

Main functions

initOs

Most of this is what is specified in the homework. The ideas are, if CurProcid is 0xFFFF it

```
165 void initOs() {  
166     mem[Cur_Proc_ID] = 0xFFFF;  
167     mem[Proc_Count] = 0x0000;  
168     mem[OS_STATUS] = 0x0000;  
169     mem[OS_FREE_BITMAP] = 0x1FFF; //first three pages are full  
170     mem[OS_FREE_BITMAP+1] = 0xFFFF;  
171     return;  
172 }
```

means no active process. There are 0 processes yet so the count is zero. Os_Status aims to signal the memory filling, that is not the case at the start. There are two

freebitmaps holding 32 pages the system has. First 3 is full because they are reserved for system. Hence first free bitmap is 0001 1111 1111 1111.

CreateProc

Aim is to create process out of the given code and heap parts of the process.

```

174 int createProc(char *fname, char *hname) {
175     if (mem[OS_STATUS] & VALID_MASK){
176         //os_stat is not zero, it is full
177         printf("The OS memory region is full. Cannot create a new PCB.\n");
178         return -1;
179     }
180     uint16_t pid = mem[Proc_Count];
181     mem[Proc_Count] += 1;
182     if (mem[Proc_Count] == 1361)
183     {
184         mem[OS_STATUS] = 0xFFFF;
185     }
186     uint16_t self_pcb = getPCB(pid);
187     mem[self_pcb + PID_PCB] = pid;
188     mem[self_pcb + PC_PCB] = 0x3000; //processes start from 0x3000
189     mem[self_pcb + PTBR_PCB] = getPageTableStart(pid);
190     FILE *code_file = fopen(fname, "r");
191     if (!code_file){
192         fprintf(stderr, "Cannot read code file: %s\n", fname);
193         return -1;
194     }
195     uint16_t offsets[2];
196     for (int i = 0; i < CODE_SIZE; i++){
197         uint16_t worked = allocMem(mem[self_pcb + PTBR_PCB], i + 6, 0xFFFF, 0);
198         if (worked){
199             offsets[i] = (mem[mem[self_pcb + PTBR_PCB] + i + 6] >> PTE_SHIFT) << PTE_SHIFT;
200         }
201         else{
202             printf("Cannot create code segment.\n");
203             return -1;
204         }
205     }
206     size_t writ = getFileSize(code_file);
207     fclose(code_file);
208     ld_img(fname, offsets, writ);
209     //...

```

This function is called by the main.c, for every given input from terminal. The user enters the name of the code followed by the name of the heap. These make up a process. We first check OS_Status. And then if this is the final possible process set the Os_Status. We fill out the process control block for this new process. Open the file for code, and allocate CODE_Size(2) pages. Allocation takes the start of page table registry, virtual page number, read permission and write permission. Virtual pages start from 6 because of the way they are implemented. Processes think the first ever part of their code starts at 0x3000. If it allocates we keep record of it for the load image function. Load the image for code, and then do the same to heap.

Load Proc and AllocMem

```
225
226 void loadProc(uint16_t pid) {
227
228     mem[Cur_Proc_ID] = pid;
229
230     uint16_t pcb = getPCB(pid);
231     reg[RPC] = mem[pcb + PC_PCB];
232     reg[PTBR] = mem[pcb + PTBR_PCB];
233 }
234
235 //for example the first code starts with 0x3000 shifted 11 is 6, it is the first though. VPN should be 5 less than first 5 bits
236 //pre: ptbr is base register for page table, vpn is the first 5 bit of virt address
237 uint16_t allocMem(uint16_t ptbr, uint16_t vpn, uint16_t read, uint16_t write) {
238     for (uint16_t i = 0; i < 32; i++) {
239         uint16_t *bitmap = (i < 16) ? &mem[OS_FREE_BITMAP] : &mem[OS_FREE_BITMAP + 1];
240         uint16_t bit = 0x8000 >> (i % 16); // Mask for current page frame
241
242         // Check if the page is free
243         if (*bitmap & bit) {
244             *bitmap &= ~bit; // Mark page as used in the bitmap
245
246             // Calculate and store the PTE in the page table
247             //printf("write is %d", write);
248             uint16_t pte = (i << PTE_SHIFT) | ((read == 0xFFFF) << 1) | ((write == 0xFFFF) << 2) | VALID_MASK;
249             //printf("new pte is 0x%04X\n", pte);
250             mem[ptbr + vpn] = pte; // Write PTE directly at ptbr + vpn
251             return 1; // Allocation successful
252         }
253     }
254     fprintf(stderr, "No free page available.\n");
255     return 0; // Failure
256 }
```

Load process is simple. Since the system remembers the process by keeping two of its properties in registers(our array based pseudo register) and the id of process in OS. We set them according to the given id.

In allocMem arguments are start of the page table registry, virtual page number, permission for read and write.

First find a place in freebitmaps that keep track of it. Since there are two words that keep track of it according to which physical page we choose we take the first bitmap or second. And the bit we will use is 1000 0000 0000 0000 shifted right according to order.

We check until we find a free one. After that we create the page table entry. First five bits is the physical page number. And then we put the bits to their respective places. For example if read is permitted the variable is 0xFFFF. If it is shift 1 left and add it to the pte. Put that pte to its memory spot.

```

Edit Selection View Go Run Terminal Help  ← →  ege.balyali [SSH: 10.3.0.91]
vm.c  ×  differences2.txt
ba4 > C vm.c
315 static inline void tbrk() {
316     uint16_t vpn = reg[R0] >> PTE_SHIFT;
317     uint16_t pcb_start = getPCB(mem[Cur_Proc_ID]);
318     uint16_t pte_start = mem[pcb_start + PTBR_PCB];
319
320     if ((mem[pte_start + vpn] & VALID_MASK) == 0) {
321         //not allocated
322         printf("Heap increase requested by process %d.\n", mem[Cur_Proc_ID]);
323         //printf("regR0 is %d", reg[R0]);
324         allocMem(mem[pcb_start + PTBR_PCB], vpn, (reg[R0] & 0x0002 ? UINT16_MAX : 0), (reg[R0] & 0x0004 ? UINT16_MAX : 0));
325     }
326     else{
327         freeMem(vpn, mem[pcb_start + PTBR_PCB]);
328     }
329 }
```

Tbrk is called to increase or decrease the heap. How do we know which one? The r0 will keep the page specified by the process. If it is already allocated deallocate it, else if it is new allocate it.

```

330 static inline void tyld() {
331     uint16_t pcb_start = getPCB(mem[Cur_Proc_ID]);
332     mem[pcb_start + PC_PCB] = reg[RPC];
333     mem[pcb_start + PTBR_PCB] = reg[PTBR];
334     uint16_t original_start = pcb_start;
335     pcb_start += PCB_SIZE;
336     while(pcb_start != original_start)
337     {
338         if((pcb_start + PID_PCB) == PCB_START && mem[pcb_start + PID_PCB] != 0xFFFF){
339             //printf("line331We are switching from process %d to %d.\n",mem[Cur_Proc_ID],mem[pcb_start + PID_PCB] );
340             printf("We are switching from process %d to %d.\n",mem[Cur_Proc_ID],mem[pcb_start + PID_PCB] );
341             loadProc(mem[pcb_start + PID_PCB]);
342             return;
343         }
344         if (mem[pcb_start + PID_PCB] != 0 && mem[pcb_start + PID_PCB] != 0xFFFF)
345         {
346             printf("We are switching from process %d to %d.\n",mem[Cur_Proc_ID],mem[pcb_start + PID_PCB] );
347             loadProc(mem[pcb_start + PID_PCB]);
348             return;
349         }
350         pcb_start = (pcb_start + PCB_SIZE) % (1360 * PCB_SIZE);
351         if (pcb_start == 0)
352             pcb_start = PCB_START;
353     }
354     //printf("came to line 345, it will call itself");
355     loadProc(mem[original_start + PID_PCB]);
356 }
```

Tyld is called when the process wants to give other processes a chance to run. If there is none it calls itself. We get the process control block of the current process. And increase it to access the next process control block until another one is found. At line 350 we ensure the wrap around by taking mod of it with the process control block size times amount of processes that can exist(almost 1360).

If at any point the new process control block has a working pid we load that process with loadProc. If none is found and the loop comes full circle meaning there is no other active process, vm loads the original process.

thalt

```
384 static inline void thalt() {
385     uint16_t pcb_start = getPCB(mem[Cur_Proc_ID]);
386     mem[pcb_start + PID_PCB] = 0xFFFF;
387     uint16_t pte_start = mem[pcb_start + PTBR_PCB];
388     //printf("pte location %d and pte is %d\n", (pte_start + i), mem[pte_start + i]);
389     for ( int i = 0; i < 32; i++){
390         if ((mem[pte_start + i] & VALID_MASK) == 1){
391             freeMem(i, mem[(pcb_start + PTBR_PCB)]);
392         }
393     }
394     uint16_t i = PCB_START;
395     uint16_t no_process_left = 1;
396     if (mem[PCB_START + PID_PCB] != 0xFFFF){
397         no_process_left = 0;
398     }
399     while(no_process_left && (i < (PCB_SIZE * 32 + PCB_START)))
400     {
401         if (mem[i + PID_PCB] != 0x0000 && mem[i + PID_PCB] != 0xFFFF){
402             no_process_left = 0;
403             break;
404         }
405         i += PCB_SIZE;
406     }
407     if (no_process_left)
408     {running = false;}
409     else{
410         changeProcAfterHalt();
411     }
412 }
```

Thalt is called by processes when they are finished. We changed that processes id to 1111 1111 1111 1111 to indicate it is finished. For every page it has we free the memory. Then we need to find if all processes ended. First we make no_process_left true. If any process is found that has a legit id no_process_left becomes false. If there aren't any process left vm raises the flag running=false and the vm stops. Else we change the process. I used to call tyield since the job is similar. However, this shouldn't print so another similar function is called.

```

357 static inline void changeProcAfterHalt() {
358     //there isn't wrap around the pte table
359     uint16_t pcb_start = getPCB(mem[Cur_Proc_ID]);
360     uint16_t original_start = pcb_start;
361     pcb_start += PCB_SIZE;
362
363     while(pcb_start != original_start)
364     {
365         //printf("original_start is %d pcb_start is at %d ",original_start,pcb_start);
366         if((pcb_start + PID_PCB) == PCB_START && mem[pcb_start + PID_PCB] != 0xFFFF)
367         {
368             loadProc(mem[pcb_start + PID_PCB]);
369             return;
370         }
371         if (mem[pcb_start + PID_PCB] != 0 && mem[pcb_start + PID_PCB] != 0xFFFF)
372         {
373             loadProc(mem[pcb_start + PID_PCB]);
374             return;
375         }
376         pcb_start = (pcb_start + PCB_SIZE) % (1360 * PCB_SIZE);
377         if (pcb_start == 0)
378             pcb_start = PCB_START;
379     }
380     //printf("came to line 345, it will call itself");
381     loadProc(mem[original_start + PID_PCB]);
382 }

```

This is the function called by thalt when there are other processes The idea is the same as tyield.

Mr and Mw

```
415
416 static inline uint16_t mr(uint16_t address) {
417     //printf("mr calls translate with %d", address);
418     uint16_t phys_address = translateAddress(address);
419     return mem[phys_address];
420 }
421
422 static inline void mw(uint16_t address, uint16_t val) {
423     //printf("mw calls translate with %d", address);
424
425     if ((getPTE(address) & WRITE_MASK) == 0)
426     {
427         printf("Cannot write to a read-only page.\n");
428         exit(1);
429     }
430     uint16_t phys_address = translateAddress(address);
431     mem[phys_address] = val;
432 }
433
434 // YOUR CODE ENDS HERE
```

Memory read translates virtual address and returns the intended information in physical memory.

Memory write is similar with the difference of checking if it is a writable page.